

# Asynchronous Streaming, and Visual Profiling with CUDA C/C++



The CUDA toolkit ships with the **Nsight Systems**, a powerful GUI application to support the development of accelerated CUDA applications. Nsight Systems generates a graphical timeline of an accelerated application, with detailed information about CUDA API calls, kernel execution, memory activity, and the use of **CUDA streams**.

In this lab, you will be using the Nsight Systems timeline to guide you in optimizing accelerated applications. Additionally, you will learn some intermediate CUDA programming techniques to support your work: **unmanaged memory allocation and migration**; **pinning**, or **page-locking** host memory; and **non-default concurrent CUDA streams**.

At the end of this lab, you will be presented with an assessment, to accelerate and optimize a simple n-body particle simulator, which will allow you to demonstrate the skills you have developed during this course. Those of you who are able to accelerate the simulator while maintaining its correctness, will be granted a certification as proof of your competency.

---

## Prerequisites

To get the most out of this lab you should already be able to:

- Write, compile, and run C/C++ programs that both call CPU functions and launch GPU kernels.
- Control parallel thread hierarchy using execution configuration.
- Refactor serial loops to execute their iterations in parallel on a GPU.
- Allocate and free CUDA Unified Memory.
- Understand the behavior of Unified Memory with regard to page faulting and data migrations.
- Use asynchronous memory prefetching to reduce page faults and data migrations.

## Objectives

By the time you complete this lab you will be able to:

- Use **Nsight Systems** to visually profile the timeline of GPU-accelerated CUDA applications.
  - Use Nsight Systems to identify, and exploit, optimization opportunities in GPU-accelerated CUDA applications.
  - Utilize CUDA streams for concurrent kernel execution in accelerated applications.
  - **(Optional Advanced Content)** Use manual device memory allocation, including allocating pinned memory, in order to asynchronously transfer data in concurrent CUDA streams.
- 

## Running Nsight Systems

For this interactive lab environment, we have set up a remote desktop you can access from your browser, where you will be able to launch and use Nsight Systems.

You will begin by creating a report file for an already-existing vector addition program, after which you will be walked through a series of steps to open this report file in Nsight Systems, and to make the visual experience nice.

### Generate Report File

[01-vector-add.cu](#) (<----- click on these links to source files to edit them in the browser) contains a working, accelerated, vector addition application. Use the code execution cell directly below (you can execute it, and any of the code execution cells in this lab by **CTRL** + clicking it) to compile and run it. You should see a message printed that indicates it was successful.

```
In [5]: !nvcc -o vector-add-no-prefetch 01-vector-add/01-vector-add.cu -run
```

Success! All values calculated correctly.

Next, use `nsys profile --stats=true` to create a report file that you will be able to open in the Nsight Systems visual profiler. Here we use the `-o` flag to give the report file a memorable name:

```
In [6]: !nsys profile --stats=true -o vector-add-no-prefetch-report ./vector-add-no-prefetch
```

Warning: LBR backtrace method is not supported on this platform. DWARF backtrace method will be used.  
 WARNING: The command line includes a target application therefore the CPU context-switch scope has been set to process-tree.  
 Collecting data...  
 Success! All values calculated correctly.  
 Processing events...  
 Saving temporary "/tmp/nsys-report-f265-10a1-ab68-76c3.qdstrm" file to disk...

Creating final output files...  
 Processing [=====100%]  
 Saved report file to "/tmp/nsys-report-f265-10a1-ab68-76c3.qdrep"  
 Exporting 10077 events: [=====100%]

Exported successfully to  
 /tmp/nsys-report-f265-10a1-ab68-76c3.sqlite

#### CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
60.7	148298522	3	49432840.7	13951	148254781	cudaMallocManaged
31.0	75803281	1	75803281.0	75803281	75803281	cudaDeviceSynchronize
8.3	20250826	3	6750275.3	5993438	7840882	cudaFree
0.0	51681	1	51681.0	51681	51681	cudaLaunchKernel

#### CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	75794123	1	75794123.0	75794123	75794123	addVectorsInto(float*, float*, float*, int)

#### CUDA Memory Operation Statistics (by time):

Time(%)	Total Time (ns)	Operations	Average	Minimum	Maximum	Operation
80.0	43237264	8203	5270.9	1822	79360	[CUDA Unified Memory memcpy HtoD]
20.0	10837032	768	14110.7	1375	80416	[CUDA Unified Memory memcpy DtoH]

# CUDA Memory Operation Statistics (by size in KiB):

Total	Operations	Average	Minimum	Maximum	Operation
393216.000	8203	47.936	4.000	1012.000	[CUDA Unified Memory memcpy HtoD]
131072.000	768	170.667	4.000	1020.000	[CUDA Unified Memory memcpy DtoH]

# Operating System Runtime API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
86.4	1777406599	98	18136802.0	24720	100140281	poll
9.2	189426307	86	2202631.5	14860	20414149	sem_timedwait
2.3	46375028	621	74678.0	1030	12386334	ioctl
1.0	20165135	23	876745.0	1130	7806771	mmap
1.0	19868203	23	863834.9	1140	19780418	fopen
0.1	1935496	64	30242.1	2480	567880	mmap64
0.0	811893	76	10682.8	3761	45091	open64
0.0	140994	11	12817.6	8310	25521	write
0.0	139843	4	34960.8	27721	45440	pthread_create
0.0	48211	1	48211.0	48211	48211	fgets
0.0	36701	12	3058.4	1380	5160	munmap
0.0	26631	5	5326.2	2470	8320	open
0.0	15090	3	5030.0	1170	9880	fread
0.0	13640	3	4546.7	1390	7610	fgetc
0.0	11620	6	1936.7	1070	3040	read
0.0	11090	6	1848.3	1040	3570	fclose
0.0	11080	2	5540.0	3060	8020	socket
0.0	7820	1	7820.0	7820	7820	connect
0.0	7290	3	2430.0	1060	5080	fcntl
0.0	7190	1	7190.0	7190	7190	pipe2
0.0	3700	1	3700.0	3700	3700	sem_wait
0.0	1770	1	1770.0	1770	1770	bind

Report file moved to "/dli/task/vector-add-no-prefetch-report.qdrep"

Report file moved to "/dli/task/vector-add-no-prefetch-report.sqlite"

## Open the Remote Desktop


Run the next cell to generate a link to the remote desktop. Then, read the instructions that follow in the notebook.

```
In [7]: %%js
var port = ((window.location.port == 80) ? "" : (":"+window.location.port));
var url = 'http://' + window.location.hostname + port + '/nsight/vnc.html?resize=scale';
let a = document.createElement('a');
a.setAttribute('href', url)
a.setAttribute('target', '_blank')
a.innerText = 'Click to open remote desktop'
element.append(a);
```

After clicking the *Connect* button you will be asked for a password, which is `nvidia`.

## Open Nsight Systems

To open Nsight Systems, double-click the "NVIDIA Nsight Systems" icon on the remote desktop.



## Enable Usage Reporting

When prompted, click "Yes" to enable usage reporting:



## Select GPU Rows on Top

When prompted, select *GPU Rows on Top* and then click *Okay*.




## Open the Report File

Open this report file by visiting *File -> Open* from the Nsight Systems menu and select `vector-add-no-prefetch-report.qdrep`:



## Ignore Warnings/Errors

You can close and ignore any warnings or errors you see, which are just a result of our particular remote desktop environment:



## Make More Room for the Timelines

To make your experience nicer, full-screen the profiler, close the *Project Explorer* and hide the *Events View*:




Your screen should now look like this:



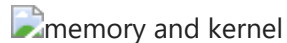
## Expand the CUDA Unified Memory Timelines

Next, expand the *CUDA* -> *Unified memory* and *Context* timelines, and close the *Threads* timelines:



## Observe Many Memory Transfers

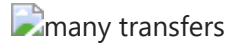
From a glance you can see that your application is taking about 1 second to run, and that also, during the time when the `addVectorsInto` kernel is running, that there is a lot of UM memory activity:



Zoom into the memory timelines to see more clearly all the small memory transfers being caused by the on-demand memory page faults. A couple tips:

1. You can zoom in and out at any point of the timeline by holding `CTRL` while scrolling your mouse/trackpad
2. You can zoom into any section by click + dragging a rectangle around it, and then selecting *Zoom in*

Here's an example of zooming in to see the many small memory transfers:



---

## Comparing Code Refactors Iteratively with Nsight Systems

Now that you have Nsight Systems up and running and are comfortable moving around the timelines, you will be profiling a series of programs that were iteratively improved using techniques already familiar to you. Each time you profile, information in the timeline will give information supporting how you should next modify your code. Doing this will further increase your understanding of how various CUDA programming techniques affect application performance.

### Exercise: Compare the Timelines of Prefetching vs. Non-Prefetching

[01-vector-add-prefetch-solution.cu](#) refactors the vector addition application from above so that the 3 vectors needed by its `addVectorsInto` kernel are asynchronously prefetched to the active GPU device prior to launching the kernel (using `cudaMemPrefetchAsync`). Open the source code and identify where in the application these changes were made.

After reviewing the changes, compile and run the refactored application using the code execution cell directly below. You should see its success message printed.

```
In [8]: !nvcc -o vector-add-prefetch 01-vector-add/solutions/01-vector-add-prefetch-solution.cu -run
```

Success! All values calculated correctly.

Now create a report file for this version of the application:

```
In [9]: !nsys profile --stats=true -o vector-add-prefetch-report ./vector-add-prefetch
```

Warning: LBR backtrace method is not supported on this platform. DWARF backtrace method will be used.  
WARNING: The command line includes a target application therefore the CPU context-switch scope has been set to process-tree.

Collecting data...

Success! All values calculated correctly.

Processing events...

Saving temporary "/tmp/nsys-report-a3af-bbd9-98f1-faa4.qdstrm" file to disk...

Creating final output files...

Processing [=====100%]

Saved report file to "/tmp/nsys-report-a3af-bbd9-98f1-faa4.qdrep"

Exporting 2082 events: [=====100%]

Exported successfully to

/tmp/nsys-report-a3af-bbd9-98f1-faa4.sqlite

#### CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
74.0	152144398	3	50714799.3	14610	152099457	cudaMallocManaged
9.6	19673682	3	6557894.0	136683	10809113	cudaMemPrefetchAsync
9.6	19646641	3	6548880.3	5849574	7438613	cudaFree
6.8	14056831	1	14056831.0	14056831	14056831	cudaDeviceSynchronize
0.0	32100	1	32100.0	32100	32100	cudaLaunchKernel

#### CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	849855	1	849855.0	849855	849855	addVectorsInto(float*, float*, float*, int)

#### CUDA Memory Operation Statistics (by time):

Time(%)	Total Time (ns)	Operations	Average	Minimum	Maximum	Operation
73.9	30536557	192	159044.6	158720	159487	[CUDA Unified Memory memcpy HtoD]
26.1	10807623	768	14072.4	1374	78559	[CUDA Unified Memory memcpy DtoH]



#### CUDA Memory Operation Statistics (by size in KiB):

Total	Operations	Average	Minimum	Maximum	Operation
393216.000	192	2048.000	2048.000	2048.000	[CUDA Unified Memory memcpy HtoD]
131072.000	768	170.667	4.000	1020.000	[CUDA Unified Memory memcpy DtoH]

#### Operating System Runtime API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
85.1	1727067489	95	18179657.8	24950	100133139	poll
9.0	182675555	81	2255253.8	14750	20415335	sem_timedwait
3.8	77151711	630	122463.0	1040	10719812	ioctl
1.0	19525661	23	848941.8	1120	7389342	mmap
0.9	18299666	21	871412.7	1300	18182735	fopen
0.1	2045267	3	681755.7	14580	1983666	sem_wait
0.1	1959180	64	30612.2	2680	574220	mmap64
0.0	803016	76	10566.0	5130	45471	open64
0.0	202644	5	40528.8	30231	60261	pthread_create
0.0	164580	12	13715.0	8590	28150	write
0.0	57101	1	57101.0	57101	57101	fgets
0.0	50992	13	3922.5	1380	15530	munmap
0.0	32410	5	6482.0	3570	11580	open
0.0	24581	12	2048.4	1060	3390	fclose
0.0	16531	4	4132.8	1020	9540	fgetc
0.0	16380	8	2047.5	1060	3180	read
0.0	12740	2	6370.0	3810	8930	socket
0.0	9100	3	3033.3	1120	5010	fread
0.0	9090	1	9090.0	9090	9090	connect
0.0	7480	1	7480.0	7480	7480	pipe2
0.0	5440	2	2720.0	1080	4360	fcntl
0.0	1960	1	1960.0	1960	1960	bind
0.0	1030	1	1030.0	1030	1030	listen

Report file moved to "/dli/task/vector-add-prefetch-report.qdrep"

Report file moved to "/dli/task/vector-add-prefetch-report.sqlite"

Open the report in Nsight Systems, leaving the previous report open for comparison.

- How does the execution time compare to that of the `addVectorsInto` kernel prior to adding asynchronous prefetching?

- Locate `cudaMemPrefetchAsync` in the *CUDA API* section of the timeline.
- How have the memory transfers changed?

## Exercise: Profile Refactor with Launch Init in Kernel

In the previous iteration of the vector addition application, the vector data is being initialized on the CPU, and therefore needs to be migrated to the GPU before the `addVectorsInto` kernel can operate on it.

The next iteration of the application, [01-init-kernel-solution.cu](#), the application has been refactored to initialize the data in parallel on the GPU.

Since the initialization now takes place on the GPU, prefetching has been done prior to initialization, rather than prior to the vector addition work. Review the source code to identify where these changes have been made.

After reviewing the changes, compile and run the refactored application using the code execution cell directly below. You should see its success message printed.

```
In [10]: !nvcc -o init-kernel 02-init-kernel/solutions/01-init-kernel-solution.cu -run  
Success! All values calculated correctly.
```

Now create a report file for this version of the application:

```
In [11]: !nsys profile --stats=true -o init-kernel-report ./init-kernel
```

Warning: LBR backtrace method is not supported on this platform. DWARF backtrace method will be used.  
WARNING: The command line includes a target application therefore the CPU context-switch scope has been set to process-tree.

Collecting data...

Success! All values calculated correctly.

Processing events...

Saving temporary "/tmp/nsys-report-d8ba-a084-fb9b-235f.qdstrm" file to disk...

Creating final output files...

Processing [=====100%]

Saved report file to "/tmp/nsys-report-d8ba-a084-fb9b-235f.qdrep"

Exporting 1769 events: [=====100%]

Exported successfully to

/tmp/nsys-report-d8ba-a084-fb9b-235f.sqlite

#### CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
88.6	162006472	3	54002157.3	20910	161905280	cudaMallocManaged
9.2	16911395	3	5637131.7	4178913	8445537	cudaFree
1.2	2166428	3	722142.7	696513	762653	cudaMemPrefetchAsync
0.9	1684859	1	1684859.0	1684859	1684859	cudaDeviceSynchronize
0.0	59541	4	14885.3	5900	37241	cudaLaunchKernel

#### CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
50.2	856575	1	856575.0	856575	856575	addVectorsInto(float*, float*, float*, int)
49.8	848478	3	282826.0	280128	284447	initWith(float, float*, int)

#### CUDA Memory Operation Statistics (by time):

Time(%)	Total Time (ns)	Operations	Average	Minimum	Maximum	Operation
100.0	10865936	768	14148.4	1439	80704	[CUDA Unified Memory memcpy DtoH]

# CUDA Memory Operation Statistics (by size in KiB):

Total	Operations	Average	Minimum	Maximum	Operation
-----	-----	-----	-----	-----	-----
131072.000	768	170.667	4.000	1020.000	[CUDA Unified Memory memcpy DtoH]

## Operating System Runtime API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
-----	-----	-----	-----	-----	-----	-----
75.2	481447946	32	15045248.3	30470	100139254	poll
11.1	70866827	27	2624697.3	15871	20551968	sem_timedwait
7.5	48115636	628	76617.3	1040	9147949	ioctl
3.0	19433847	22	883356.7	1000	19324526	fopen
2.6	16822413	24	700933.9	1010	8384286	mmap
0.3	2009097	64	31392.1	2500	564999	mmap64
0.1	807761	76	10628.4	3390	29821	open64
0.0	163121	4	40780.3	30210	61641	pthread_create
0.0	157142	11	14285.6	11040	19211	write
0.0	48481	1	48481.0	48481	48481	fgets
0.0	42830	12	3569.2	1790	7870	munmap
0.0	28952	5	5790.4	2420	8501	open
0.0	24250	4	6062.5	1820	15040	fread
0.0	17170	7	2452.9	1020	5690	fclose
0.0	15240	4	3810.0	1170	7550	fgetc
0.0	13541	2	6770.5	5140	8401	socket
0.0	13211	8	1651.4	1060	3670	read
0.0	8860	1	8860.0	8860	8860	connect
0.0	7480	1	7480.0	7480	7480	pipe2
0.0	7040	2	3520.0	1040	6000	fcntl
0.0	2290	1	2290.0	2290	2290	bind
0.0	1190	1	1190.0	1190	1190	listen

Report file moved to "/dli/task/init-kernel-report.qdrep"

Report file moved to "/dli/task/init-kernel-report.sqlite"

Open the new report file in Nsight Systems and do the following:

- Compare the application and `addVectorsInto` run times to the previous version of the application, how did they change?

- Look at the *Kernels* section of the timeline. Which of the two kernels ( `addVectorsInto` and the initialization kernel) is taking up the majority of the time on the GPU?
- Which of the following does your application contain?
  - Data Migration (HtoD)
  - Data Migration (DtoH)

## Exercise: Profile Refactor with Asynchronous Prefetch Back to the Host

Currently, the vector addition application verifies the work of the vector addition kernel on the host. The next refactor of the application, [01-prefetch-check-solution.cu](#), asynchronously prefetches the data back to the host for verification.

After reviewing the changes, compile and run the refactored application using the code execution cell directly below. You should see its success message printed.

```
In [12]: !nvcc -o prefetch-to-host 04-prefetch-check/solutions/01-prefetch-check-solution.cu -run
```

Success! All values calculated correctly.

Now create a report file for this version of the application:

```
In [13]: !nsys profile --stats=true -o prefetch-to-host-report ./prefetch-to-host
```

Warning: LBR backtrace method is not supported on this platform. DWARF backtrace method will be used.  
WARNING: The command line includes a target application therefore the CPU context-switch scope has been set to process-tree.

Collecting data...

Success! All values calculated correctly.

Processing events...

Saving temporary "/tmp/nsys-report-7cad-1f62-2474-6f32.qdstrm" file to disk...

Creating final output files...

Processing [=====100%]

Saved report file to "/tmp/nsys-report-7cad-1f62-2474-6f32.qdrep"

Exporting 1057 events: [=====100%]

Exported successfully to

/tmp/nsys-report-7cad-1f62-2474-6f32.sqlite

#### CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
78.2	152915703	3	50971901.0	14301	152841041	cudaMallocManaged
13.0	25351408	4	6337852.0	513809	23729960	cudaMemPrefetchAsync
8.0	15610730	3	5203576.7	4143812	7150374	cudaFree
0.9	1688789	1	1688789.0	1688789	1688789	cudaDeviceSynchronize
0.0	48841	4	12210.3	4651	31710	cudaLaunchKernel

#### CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
50.4	858654	1	858654.0	858654	858654	addVectorsInto(float*, float*, float*, int)
49.6	844350	3	281450.0	279135	282751	initWith(float, float*, int)

#### CUDA Memory Operation Statistics (by time):

Time(%)	Total Time (ns)	Operations	Average	Minimum	Maximum	Operation
100.0	9991305	64	156114.1	155935	156799	[CUDA Unified Memory memcpy DtoH]

# CUDA Memory Operation Statistics (by size in KiB):

Total	Operations	Average	Minimum	Maximum	Operation
-----	-----	-----	-----	-----	-----
131072.000	64	2048.000	2048.000	2048.000	[CUDA Unified Memory memcpy DtoH]

## Operating System Runtime API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
-----	-----	-----	-----	-----	-----	-----
68.2	341994799	27	12666474.0	20310	100129950	poll
14.9	74909308	629	119092.7	1000	23685749	ioctl
13.1	65860455	23	2863498.0	14431	20454453	sem_timedwait
3.1	15520750	22	705488.6	1100	7120193	mmap
0.4	1986120	64	31033.1	2940	583811	mmap64
0.2	789856	76	10392.8	3220	38071	open64
0.0	152551	4	38137.8	29050	51911	pthread_create
0.0	142602	11	12963.8	8440	21060	write
0.0	102110	22	4641.4	1130	19620	fopen
0.0	48640	1	48640.0	48640	48640	fgets
0.0	48130	11	4375.5	1300	16510	munmap
0.0	27590	5	5518.0	2460	8480	open
0.0	17930	4	4482.5	1390	8000	fgetc
0.0	12850	6	2141.7	1130	3780	read
0.0	12680	1	12680.0	12680	12680	sem_wait
0.0	10990	4	2747.5	1700	3760	fread
0.0	10880	7	1554.3	1030	2920	fclose
0.0	10680	2	5340.0	3530	7150	socket
0.0	9590	5	1918.0	1080	4520	fcntl
0.0	6910	1	6910.0	6910	6910	pipe2
0.0	6570	1	6570.0	6570	6570	connect
0.0	1660	1	1660.0	1660	1660	bind
0.0	1000	1	1000.0	1000	1000	listen

Report file moved to "/dli/task/prefetch-to-host-report.qdrep"

Report file moved to "/dli/task/prefetch-to-host-report.sqlite"

Open this report file in Nsight Systems, and do the following:

- Use the *Unified Memory* section of the timeline to compare and contrast the *Data Migration (DtoH)* events before and after adding prefetching back to the CPU.
- 

## Concurrent CUDA Streams

You are now going to learn about a new concept, **CUDA Streams**. After an introduction to them, you will return to using Nsight Systems to better evaluate their impact on your application's performance.

The following video presents upcoming material visually, at a high level. Click watch it before moving on to more detailed coverage of their topics in following sections.

```
In [14]: from IPython.display import HTML

video_url = "https://d36m44n9vdbmda.cloudfront.net/assets/s-ac-04-v1/task3/NVVP-Streams-1.mp4"

video_html = f"""
<video controls width="640" height="360">
    <source src="{video_url}" type="video/mp4">
    Your browser does not support the video tag.
</video>
"""

display(HTML(video_html))
```





In CUDA programming, a **stream** is a series of commands that execute in order. In CUDA applications, kernel execution, as well as some memory transfers, occur within CUDA streams. Up until this point in time, you have not been interacting explicitly with CUDA streams, but in fact, your CUDA code has been executing its kernels inside of a stream called *the default stream*.

CUDA programmers can create and utilize non-default CUDA streams in addition to the default stream, and in doing so, perform multiple operations, such as executing multiple kernels, concurrently, in different streams. Using multiple streams can add an additional layer of parallelization to your accelerated applications, and offers many more opportunities for application optimization.

## Rules Governing the Behavior of CUDA Streams

There are a few rules, concerning the behavior of CUDA streams, that should be learned in order to utilize them effectively:

- Operations within a given stream occur in order.
- Operations in different non-default streams are not guaranteed to operate in any specific order relative to each other.
- The default stream is blocking and will both wait for all other streams to complete before running, and, will block other streams from running until it completes.

## Creating, Utilizing, and Destroying Non-Default CUDA Streams

The following code snippet demonstrates how to create, utilize, and destroy a non-default CUDA stream. You will note, that to launch a CUDA kernel in a non-default CUDA stream, the stream must be passed as the optional 4th argument of the execution configuration. Up until now you have only utilized the first 2 arguments of the execution configuration:

```
cudaStream_t stream;          // CUDA streams are of type `cudaStream_t`.
cudaStreamCreate(&stream);    // Note that a pointer must be passed to `cudaCreateStream`.

someKernel<<<number_of_blocks, threads_per_block, 0, stream>>>(); // `stream` is passed as 4th EC
argument.

cudaStreamDestroy(stream);    // Note that a value, not a pointer, is passed to `cudaDestroyStream`.
```

Outside the scope of this lab, but worth mentioning, is the optional 3rd argument of the execution configuration. This argument allows programmers to supply the number of bytes in **shared memory** (an advanced topic that will not be covered presently) to be dynamically allocated per block for this kernel launch. The default number of bytes allocated to shared memory per block is `0`, and for the remainder of the lab, you will be passing `0` as this value, in order to expose the 4th argument, which is of immediate interest:

### Exercise: Predict Default Stream Behavior

The `01-print-numbers` application has a very simple `printNumber` kernel which accepts an integer and prints it. The kernel is only being executed with a single thread inside a single block. However, it is being executed 5 times, using a for-loop, and passing each launch the number of the for-loop's iteration.

Compile and run `01-print-numbers` using the code execution block below. You should see the numbers `0` through `4` printed.

```
In [15]: !nvcc -o print-numbers 05-stream-intro/01-print-numbers.cu -run
```

```
0
1
2
3
4
```

Knowing that by default kernels are executed in the default stream, would you expect that the 5 launches of the `print-numbers` program executed serially, or in parallel? You should be able to mention two features of the default stream to support your answer. Create a report file in the cell below and open it in Nsight Systems to confirm your answer.

```
In [16]: !nsys profile --stats=true -o print-numbers-report ./print-numbers
```

Warning: LBR backtrace method is not supported on this platform. DWARF backtrace method will be used.  
WARNING: The command line includes a target application therefore the CPU context-switch scope has been set to process-tree.

Collecting data...

0  
1  
2  
3  
4

Processing events...

Saving temporary "/tmp/nsys-report-36c2-2d80-4fe1-80d9.qdstrm" file to disk...

Creating final output files...

Processing [=====100%]

Saved report file to "/tmp/nsys-report-36c2-2d80-4fe1-80d9.qdrep"

Exporting 941 events: [=====100%]

Exported successfully to

/tmp/nsys-report-36c2-2d80-4fe1-80d9.sqlite

#### CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.9	136181700	5	27236340.0	4520	136158989	cudaLaunchKernel
0.1	156602	1	156602.0	156602	156602	cudaDeviceSynchronize

#### CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	148128	5	29625.6	28128	33728	printNumber(int)

#### Operating System Runtime API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
63.5	134499923	13	10346147.9	18841	45072014	poll
25.3	53446117	625	85513.8	1000	9676686	ioctl
9.5	20042816	21	954419.8	1240	19960393	fopen

0.9	1914382	64	29912.2	2330	573389	mmap64
0.4	775851	76	10208.6	3240	34130	open64
0.1	282513	10	28251.3	14710	61231	sem_timedwait
0.1	156783	12	13065.3	8460	23501	write
0.1	148132	4	37033.0	33180	46450	pthread_create
0.1	141422	17	8318.9	1410	33000	mmap
0.0	72211	8	9026.4	1850	50501	munmap
0.0	48681	1	48681.0	48681	48681	fgets
0.0	36991	5	7398.2	2400	21511	open
0.0	16500	3	5500.0	1180	10770	fread
0.0	13150	3	4383.3	1310	7110	fgetc
0.0	12640	1	12640.0	12640	12640	sem_wait
0.0	12251	2	6125.5	3800	8451	socket
0.0	11601	7	1657.3	1081	2330	read
0.0	11500	6	1916.7	1040	3150	fclose
0.0	8000	1	8000.0	8000	8000	connect
0.0	7280	3	2426.7	1010	4990	fcntl
0.0	7160	1	7160.0	7160	7160	pipe2
0.0	1620	1	1620.0	1620	1620	bind

Report file moved to "/dli/task/print-numbers-report.qdrep"

Report file moved to "/dli/task/print-numbers-report.sqlite"

## Exercise: Implement Concurrent CUDA Streams

Both because all 5 kernel launches occurred in the same stream, you should not be surprised to have seen that the 5 kernels executed serially. Additionally you could make the case that because the default stream is blocking, each launch of the kernel would wait to complete before the next launch, and this is also true.

Refactor [01-print-numbers](#) so that each kernel launch occurs in its own non-default stream. Be sure to destroy the streams you create after they are no longer needed. Compile and run the refactored code with the code execution cell directly below. You should still see the numbers `0` through `4` printed, though not necessarily in ascending order. Refer to [the solution](#) if you get stuck.

```
In [17]: !nvcc -o print-numbers-in-streams 05-stream-intro/01-print-numbers.cu -run
```

```
0
1
2
3
4
```

Now that you are using 5 different non-default streams for each of the 5 kernel launches, do you expect that they will run serially or in parallel? In addition to what you now know about streams, take into account how trivial the `printNumber` kernel is, meaning, even if you predict parallel runs, will the speed at which one kernel will complete allow for complete overlap?

After hypothesizing, open a new report file in Nsight Systems to view its actual behavior. You should notice that now, there are additional rows in the *CUDA* section for each of the non-default streams you created:

```
In [18]: !nsys profile --stats=true -o print-numbers-in-streams-report print-numbers-in-streams
```

Warning: LBR backtrace method is not supported on this platform. DWARF backtrace method will be used.  
WARNING: The command line includes a target application therefore the CPU context-switch scope has been set to process-tree.

Collecting data...

0

1

2

3

4

Processing events...

Saving temporary "/tmp/nsys-report-977f-b500-1785-02d4.qdstrm" file to disk...

Creating final output files...

Processing [=====100%]

Saved report file to "/tmp/nsys-report-977f-b500-1785-02d4.qdrep"

Exporting 965 events: [=====100%]

Exported successfully to

/tmp/nsys-report-977f-b500-1785-02d4.sqlite

#### CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.1	147694977	5	29538995.4	2010	147683527	cudaStreamCreate
0.9	1273422	5	254684.4	5670	1245131	cudaLaunchKernel
0.0	69801	1	69801.0	69801	69801	cudaDeviceSynchronize
0.0	23101	5	4620.2	2350	7371	cudaStreamDestroy

#### CUDA Kernel Statistics:


Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
100.0	225854	5	45170.8	34048	58592	printNumber(int)

#### Operating System Runtime API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
63.3	147143478	13	11318729.1	32710	46507120	poll

27.3	63486696	628	101093.5	1040	9829697	ioctl
7.7	17820573	23	774807.5	1430	17712751	fopen
0.9	1992916	64	31139.3	2570	571410	mmap64
0.3	801351	76	10544.1	4591	33061	open64
0.2	574259	10	57425.9	23170	231264	sem_timedwait
0.1	170884	12	14240.3	11430	18320	write
0.1	147832	4	36958.0	30540	48441	pthread_create
0.1	124582	17	7328.4	1150	34131	mmap
0.0	58041	1	58041.0	58041	58041	fgets
0.0	39951	5	7990.2	1290	21131	fgetc
0.0	31560	5	6312.0	3570	9440	open
0.0	23822	7	3403.1	2260	4070	munmap
0.0	19060	11	1732.7	1030	2920	fclose
0.0	18960	13	1458.5	1010	3270	read
0.0	17150	2	8575.0	6150	11000	socket
0.0	14551	1	14551.0	14551	14551	pipe2
0.0	11830	1	11830.0	11830	11830	connect
0.0	5240	2	2620.0	2320	2920	fread
0.0	4290	1	4290.0	4290	4290	fcntl
0.0	3060	1	3060.0	3060	3060	bind
0.0	1240	1	1240.0	1240	1240	listen

Report file moved to "/dli/task/print-numbers-in-streams-report.qdrep"  
 Report file moved to "/dli/task/print-numbers-in-streams-report.sqlite"

 streams print

## Exercise: Use Streams for Concurrent Data Initialization Kernels

The vector addition application you have been working with, [01-prefetch-check-solution.cu](#), currently launches an initialization kernel 3 times - once each for each of the 3 vectors needing initialization for the `vectorAdd` kernel. Refactor it to launch each of the 3 initialization kernel launches in their own non-default stream. You should still see the success message print when compiling and running with the code execution cell below. Refer to [the solution](#) if you get stuck.

```
In [19]: !nvcc -o init-in-streams 04-prefetch-check/solutions/01-prefetch-check-solution.cu -run
```

Success! All values calculated correctly.

Open a report in Nsight Systems to confirm that your 3 initialization kernel launches are running in their own non-default streams, with some degree of concurrent overlap.



```
In [20]: !nsys profile --stats=true -o init-in-streams-report ./init-in-streams
```

Warning: LBR backtrace method is not supported on this platform. DWARF backtrace method will be used.  
 WARNING: The command line includes a target application therefore the CPU context-switch scope has been set to process-tree.  
 Collecting data...  
 Success! All values calculated correctly.  
 Processing events...  
 Saving temporary "/tmp/nsys-report-0157-ab65-f1a5-23d6.qdstrm" file to disk...

Creating final output files...  
 Processing [=====100%]  
 Saved report file to "/tmp/nsys-report-0157-ab65-f1a5-23d6.qdrep"  
 Exporting 1069 events: [=====100%]

Exported successfully to  
 /tmp/nsys-report-0157-ab65-f1a5-23d6.sqlite

#### CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
76.1	151833493	3	50611164.3	22320	151722872	cudaMallocManaged
14.9	29726116	4	7431529.0	631191	27680462	cudaMemPrefetchAsync
8.0	16011217	3	5337072.3	4127099	7631087	cudaFree
0.8	1674147	1	1674147.0	1674147	1674147	cudaDeviceSynchronize
0.0	65251	4	16312.8	7650	38371	cudaLaunchKernel
0.0	60241	3	20080.3	4110	51001	cudaStreamDestroy
0.0	33860	3	11286.7	2590	28290	cudaStreamCreate

#### CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
53.3	979517	3	326505.7	298911	351487	initWith(float, float*, int)
46.7	857182	1	857182.0	857182	857182	addVectorsInto(float*, float*, float*, int)

#### CUDA Memory Operation Statistics (by time):

Time(%)	Total Time (ns)	Operations	Average	Minimum	Maximum	Operation
100.0	10023362	64	156615.0	155871	160319	[CUDA Unified Memory memcpy DtoH]

# CUDA Memory Operation Statistics (by size in KiB):

Total	Operations	Average	Minimum	Maximum	Operation
-----	-----	-----	-----	-----	-----
131072.000	64	2048.000	2048.000	2048.000	[CUDA Unified Memory memcpy DtoH]

## Operating System Runtime API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
-----	-----	-----	-----	-----	-----	-----
66.1	341446469	27	12646165.5	29550	100142691	poll
14.4	74604005	632	118044.3	1020	27613580	ioctl
12.0	61756137	23	2685049.4	18960	20632574	sem_timedwait
3.8	19690907	22	895041.2	1170	19586287	fopen
3.1	15982909	22	726495.9	1440	7592067	mmap
0.4	1935598	64	30243.7	2370	551979	mmap64
0.1	765501	76	10072.4	3620	32750	open64
0.0	149473	4	37368.3	30250	49251	pthread_create
0.0	135632	11	12330.2	8690	16781	write
0.0	50231	5	10046.2	2790	31080	open
0.0	48770	13	3751.5	1370	7920	munmap
0.0	48730	1	48730.0	48730	48730	fgets
0.0	14230	8	1778.8	1280	3590	read
0.0	14061	8	1757.6	1020	4500	fclose
0.0	13900	4	3475.0	1190	6980	fgetc
0.0	10732	2	5366.0	3061	7671	socket
0.0	10290	4	2572.5	1770	3870	fread
0.0	8930	4	2232.5	1000	5650	fcntl
0.0	7660	1	7660.0	7660	7660	connect
0.0	6500	1	6500.0	6500	6500	pipe2
0.0	1570	1	1570.0	1570	1570	bind
0.0	1210	1	1210.0	1210	1210	listen

Report file moved to "/dli/task/init-in-streams-report.qdrep"  
Report file moved to "/dli/task/init-in-streams-report.sqlite"

# Summary

At this point in the lab you are able to:

- Use the **Nsight Systems** to visually profile the timeline of GPU-accelerated CUDA applications.
- Use Nsight Systems to identify, and exploit, optimization opportunities in GPU-accelerated CUDA applications.
- Utilize CUDA streams for concurrent kernel execution in accelerated applications.

At this point in time you have a wealth of fundamental tools and techniques for accelerating CPU-only applications, and for then optimizing those accelerated applications. In the final exercise, you will have a chance to apply everything that you've learned to accelerate an [n-body](#) simulator, which predicts the individual motions of a group of objects interacting with each other gravitationally.

---

## Final Exercise: Accelerate and Optimize an N-Body Simulator

An [n-body](#) simulator predicts the individual motions of a group of objects interacting with each other gravitationally. [01-nbody.cu](#) contains a simple, though working, n-body simulator for bodies moving through 3 dimensional space.

In its current CPU-only form, this application takes about 5 seconds to run on 4096 particles, and **20 minutes** to run on 65536 particles. Your task is to GPU accelerate the program, retaining the correctness of the simulation.

## Considerations to Guide Your Work

Here are some things to consider before beginning your work:

- Especially for your first refactors, the logic of the application, the `bodyForce` function in particular, can and should remain largely unchanged: focus on accelerating it as easily as possible.
- The code base contains a for-loop inside `main` for integrating the interbody forces calculated by `bodyForce` into the positions of the bodies in the system. This integration both needs to occur after `bodyForce` runs, and, needs to complete before the next call to `bodyForce`. Keep this in mind when choosing how and where to parallelize.
- Use a **profile driven** and iterative approach.
- You are not required to add error handling to your code, but you might find it helpful, as you are responsible for your code working correctly.

## Have Fun!

Use this cell to compile the nbody simulator. Although it is initially a CPU-only application, it does accurately simulate the positions of the particles.

```
In [1]: !nvcc -std=c++11 -o nbody 09-nbody/01-nbody.cu
```

It is highly recommended you use the profiler to assist your work. Execute the following cell to generate a report file:

```
In [2]: !nsys profile --stats=true --force-overwrite=true -o nbody-report ./nbody
```

Warning: LBR backtrace method is not supported on this platform. DWARF backtrace method will be used.  
WARNING: The command line includes a target application therefore the CPU context-switch scope has been set to process-tree.

Collecting data...

4096 Bodies: average 35.810 Billion Interactions / second

Processing events...

Saving temporary "/tmp/nsys-report-5c1f-899e-ce56-15dc.qdstrm" file to disk...

Creating final output files...

Processing [=====100%]

Saved report file to "/tmp/nsys-report-5c1f-899e-ce56-15dc.qdrep"

Exporting 1057 events: [=====100%]

Exported successfully to

/tmp/nsys-report-5c1f-899e-ce56-15dc.sqlite

#### CUDA API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
99.9	3356582440	1	3356582440.0	3356582440	3356582440	cudaMallocManaged
0.1	3829524	21	182358.3	1290	1024610	cudaDeviceSynchronize
0.0	837505	20	41875.3	4330	734974	cudaLaunchKernel
0.0	65711	1	65711.0	65711	65711	cudaFree

#### CUDA Kernel Statistics:

Time(%)	Total Time (ns)	Instances	Average	Minimum	Maximum	Name
98.9	3708445	10	370844.5	297952	1021599	bodyForce(Body*, float, int)
1.1	40157	10	4015.7	3936	4416	add(Body*, float, int)

#### Operating System Runtime API Statistics:

Time(%)	Total Time (ns)	Num Calls	Average	Minimum	Maximum	Name
96.4	3345904422	48	69706342.1	23020	100155388	poll
1.8	62741156	620	101195.4	1030	13306282	ioctl
1.2	41811401	13	3216261.6	15130	20459258	sem_timedwait
0.5	16898101	24	704087.5	1731	16734037	fopen

0.1	1975477	64	30866.8	2940	545740	mmap64
0.0	796303	76	10477.7	2810	29331	open64
0.0	184172	17	10833.6	1090	67881	mmap
0.0	144743	4	36185.8	31370	44891	pthread_create
0.0	127503	11	11591.2	8440	18301	write
0.0	48431	1	48431.0	48431	48431	fgets
0.0	29630	5	5926.0	2380	8000	open
0.0	28341	9	3149.0	1470	4260	munmap
0.0	21710	3	7236.7	2010	12940	fwrite
0.0	21430	9	2381.1	1080	4650	fclose
0.0	15040	4	3760.0	1250	7010	fgetc
0.0	14810	2	7405.0	4400	10410	fopen64
0.0	14111	2	7055.5	5220	8891	fread
0.0	13160	2	6580.0	3820	9340	socket
0.0	12050	7	1721.4	1100	3260	read
0.0	10580	3	3526.7	1450	4630	fcntl
0.0	7360	1	7360.0	7360	7360	connect
0.0	5331	1	5331.0	5331	5331	pipe2
0.0	4390	1	4390.0	4390	4390	fflush
0.0	1980	1	1980.0	1980	1980	bind

Report file moved to "/dli/task/nbody-report.qdrep"  
 Report file moved to "/dli/task/nbody-report.sqlite"

Here we import a function that will run your `nbody` simulator against a various number of particles, checking for performance and accuracy.

```
In [3]: from assessment import run_assessment
```

Execute the following cell to run and assess `nbody` :

```
In [4]: run_assessment()
```

Running nbody simulator with 4096 bodies

-----

Application should run faster than 0.9s

Your application ran in: 0.1157s

Your application reports 4096 Bodies: average 44.232 Billion Interactions / second

Your results are correct

Running nbody simulator with 65536 bodies

-----

Application should run faster than 1.3s

Your application ran in: 0.2103s

Your application reports 65536 Bodies: average 430.530 Billion Interactions / second

Your results are correct

Congratulations! You passed the assessment!

See instructions below to generate a certificate, and see if you can accelerate the simulator even more!

## Generate a Certificate

If you passed the assessment, please return to the course page (shown below) and click the "ASSESS TASK" button, which will generate your certificate for the course.



## Advanced Content

The following sections, for those of you with time and interest, introduce more intermediate techniques involving some manual device memory management, and using non-default streams to overlap kernel execution and memory copies.

After learning about each of the techniques below, try to further optimize your nbody simulation using these techniques.

---

## Manual Device Memory Allocation and Copying



While `cudaMallocManaged` and `cudaMemPrefetchAsync` are performant, and greatly simplify memory migration, sometimes it can be worth it to use more manual methods for memory allocation. This is particularly true when it is known that data will only be accessed on the device or host, and the cost of migrating data can be reclaimed in exchange for the fact that no automatic on-demand migration is needed.

Additionally, using manual device memory management can allow for the use of non-default streams for overlapping data transfers with computational work. In this section you will learn some basic manual device memory allocation and copy techniques, before extending these techniques to overlap data copies with computational work.

Here are some CUDA commands for manual device memory management:

- `cudaMalloc` will allocate memory directly to the active GPU. This prevents all GPU page faults. In exchange, the pointer it returns is not available for access by host code.
- `cudaMallocHost` will allocate memory directly to the CPU. It also "pins" the memory, or page locks it, which will allow for asynchronous copying of the memory to and from a GPU. Too much pinned memory can interfere with CPU performance, so use it only with intention. Pinned memory should be freed with `cudaFreeHost`.
- `cudaMemcpy` can copy (not transfer) memory, either from host to device or from device to host.

## Manual Device Memory Management Example

Here is a snippet of code that demonstrates the use of the above CUDA API calls.

```
int *host_a, *device_a;           // Define host-specific and device-specific arrays.
cudaMalloc(&device_a, size);      // `device_a` is immediately available on the GPU.
cudaMallocHost(&host_a, size);    // `host_a` is immediately available on CPU, and is page-locked, or pinned.

initializeOnHost(host_a, N);      // No CPU page faulting since memory is already allocated on the host.

// `cudaMemcpy` takes the destination, source, size, and a CUDA-provided variable for the direction of the
// copy.
cudaMemcpy(device_a, host_a, size, cudaMemcpyHostToDevice);

kernel<<<blocks, threads, 0, someStream>>>(device_a, N);

// `cudaMemcpy` can also copy data from device to host.
cudaMemcpy(host_a, device_a, size, cudaMemcpyDeviceToHost);
```

```
verifyOnHost(host_a, N);

cudaFree(device_a);
cudaFreeHost(host_a);           // Free pinned memory like this.
```

## Exercise: Manually Allocate Host and Device Memory

The most recent iteration of the vector addition application, [01-stream-init-solution](#), is using `cudaMallocManaged` to allocate managed memory first used on the device by the initialization kernels, then on the device by the vector add kernel, and then by the host, where the memory is automatically transferred, for verification. This is a sensible approach, but it is worth experimenting with some manual device memory allocation and copying to observe its impact on the application's performance.

Refactor the [01-stream-init-solution](#) application to **not** use `cudaMallocManaged`. In order to do this you will need to do the following:

- Replace calls to `cudaMallocManaged` with `cudaMalloc`.
- Create an additional vector that will be used for verification on the host. This is required since the memory allocated with `cudaMalloc` is not available to the host. Allocate this host vector with `cudaMallocHost`.
- After the `addVectorsInto` kernel completes, use `cudaMemcpy` to copy the vector with the addition results, into the host vector you created with `cudaMallocHost`.
- Use `cudaFreeHost` to free the memory allocated with `cudaMallocHost`.

Refer to [the solution](#) if you get stuck.

```
In [21]: !nvcc -o vector-add-manual-alloc 06-stream-init/solutions/01-stream-init-solution.cu -run
```

Success! All values calculated correctly.

After completing the refactor, open a report in Nsight Systems, and use the timeline to do the following:

- Notice that there is no longer a *Unified Memory* section of the timeline.
- Comparing this timeline to that of the previous refactor, compare the run times of `cudaMalloc` in the current application vs. `cudaMallocManaged` in the previous.
- Notice how in the current application, work on the initialization kernels does not start until a later time than it did in the previous iteration. Examination of the timeline will show the difference is the time taken by `cudaMallocHost`. This clearly points out the difference between memory transfers, and memory copies. When copying memory, as you are doing presently, the data will exist

in 2 different places in the system. In the current case, the allocation of the 4th host-only vector incurs a small cost in performance, compared to only allocating 3 vectors in the previous iteration.

---

## Using Streams to Overlap Data Transfers and Code Execution


The following video presents upcoming material visually, at a high level. Click watch it before moving on to more detailed coverage of their topics in following sections.

```
In [22]: from IPython.display import HTML

video_url = "https://d36m44n9vdbmda.cloudfront.net/assets/s-ac-04-v1/task3/NVVP-Streams-3.mp4"

video_html = f"""
<video controls width="640" height="360">
    <source src="{video_url}" type="video/mp4">
    Your browser does not support the video tag.
</video>
"""

display(HTML(video_html))
```



In addition to `cudaMemcpy` is `cudaMemcpyAsync` which can asynchronously copy memory either from host to device or from device to host as long as the host memory is pinned, which can be done by allocating it with `cudaMallocHost`.

Similar to kernel execution, `cudaMemcpyAsync` is only asynchronous by default with respect to the host. It executes, by default, in the default stream and therefore is a blocking operation with regard to other CUDA operations occurring on the GPU. The `cudaMemcpyAsync` function, however, takes as an optional 5th argument, a non-default stream. By passing it a non-default stream, the memory transfer can be concurrent to other CUDA operations occurring in other non-default streams.

A common and useful pattern is to use a combination of pinned host memory, asynchronous memory copies in non-default streams, and kernel executions in non-default streams, to overlap memory transfers with kernel execution.

In the following example, rather than wait for the entire memory copy to complete before beginning work on the kernel, segments of the required data are copied and worked on, with each copy/work segment running in its own non-default stream. Using this technique, work on parts of the data can begin while memory transfers for later segments occur concurrently. Extra care must be taken when using this technique to calculate segment-specific values for the number of operations, and the offset location inside arrays, as shown here:

```

int N = 2<<24;
int size = N * sizeof(int);

int *host_array;
int *device_array;

cudaMallocHost(&host_array, size);           // Pinned host memory allocation.
cudaMalloc(&device_array, size);             // Allocation directly on the active GPU device.

initializeData(host_array, N);               // Assume this application needs to initialize on the
host.

const int numberOfSegments = 4;             // This example demonstrates slicing the work into 4
segments.
int segmentN = N / numberOfSegments;        // A value for a segment's worth of `N` is needed.
size_t segmentSize = size / numberOfSegments; // A value for a segment's worth of `size` is needed.

// For each of the 4 segments...
for (int i = 0; i < numberOfSegments; ++i)
{
    // Calculate the index where this particular segment should operate within the larger arrays.
    segmentOffset = i * segmentN;

    // Create a stream for this segment's worth of copy and work.
    cudaStream_t stream;
    cudaStreamCreate(&stream);

    // Asynchronously copy segment's worth of pinned host memory to device over non-default stream.
    cudaMemcpyAsync(&device_array[segmentOffset], // Take care to access correct location in array.
                   &host_array[segmentOffset],   // Take care to access correct location in array.
                   segmentSize,                  // Only copy a segment's worth of memory.
                   cudaMemcpyHostToDevice,
                   stream);                       // Provide optional argument for non-default stream.

    // Execute segment's worth of work over same non-default stream as memory copy.
    kernel<<<number_of_blocks, threads_per_block, 0, stream>>>(&device_array[segmentOffset], segmentN);

    // `cudaStreamDestroy` will return immediately (is non-blocking), but will not actually destroy stream
    until
    // all stream operations are complete.

```

```
    cudaStreamDestroy(stream);  
}
```

## Exercise: Overlap Kernel Execution and Memory Copy Back to Host

The most recent iteration of the vector addition application, [01-manual-malloc-solution.cu](#), is currently performing all of its vector addition work on the GPU before copying the memory back to the host for verification.

Refactor [01-manual-malloc-solution.cu](#) to perform the vector addition in 4 segments, in non-default streams, so that asynchronous memory copies can begin before waiting for all vector addition work to complete. Refer to [the solution](#) if you get stuck.

```
In [23]: !nvcc -o vector-add-manual-alloc 07-manual-malloc/solutions/01-manual-malloc-solution.cu -run
```

Success! All values calculated correctly.

After completing the refactor, open a report in Nsight Systems, and use the timeline to do the following:

- Note when the device to host memory transfers begin, is it before or after all kernel work has completed?
- Notice that the 4 memory copy segments themselves do not overlap. Even in separate non-default streams, only one memory transfer in a given direction (DtoH here) at a time can occur simultaneously. The performance gains here are in the ability to start the transfers earlier than otherwise, and it is not hard to imagine in an application where a less trivial amount of work was being done compared to a simple addition operation, that the memory copies would not only start earlier, but also overlap with kernel execution.